

CISC 372: Parallel Computing

C, part 1

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware


siegel@udel.edu

What is C?

- ▶ 1972, Dennis Ritchie, Bell Labs
- ▶ a programming language defined by an international standard
 - ▶ currently ISO/IEC 9899:2018 : Programming Language — C (“C18” or “C17”)
 - ▶ spec in docs folder on public svn repo
- ▶ characteristics
 - ▶ general purpose
 - ▶ imperative
 - ▶ static types
 - ▶ structured programming
 - ▶ lexical scopes
 - ▶ recursion
 - ▶ “low-level” (memory is a sequence of bytes, pointers; “a portable assembly language”)
- ▶ unlike modern languages, C has **unspecified** and **undefined** behavior
 - ▶ the standard leaves open many choices to the implementation
 - ▶ “unspecified”: a finite number of implementation-specific choices
 - ▶ “undefined”: anything can happen; should always be considered defects

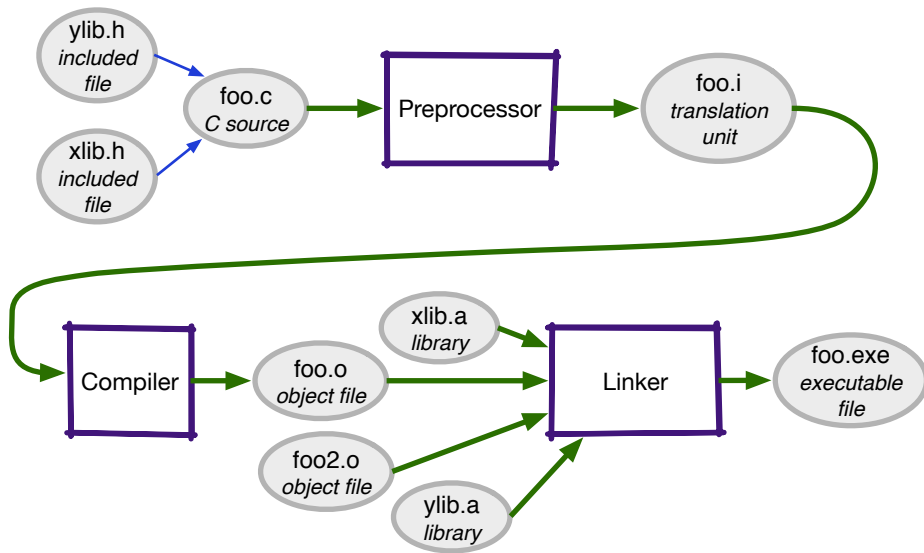
What you learn in CISC 210

What you learn in CISC 220



C++

Translation of a C program



Typical command line syntax of a C compiler

- ▶ `cc -o foo.exec foo.c`
 - ▶ preprocess, compile, and link `foo.c` creating executable `foo.exec`
 - ▶ suitable for simple programs consisting of one translation unit
- ▶ `cc -E -o foo.i foo.c`
 - ▶ preprocess only, sending output to `foo.i`
 - ▶ useful for seeing what the preprocessor is doing; debugging
- ▶ `cc -c -o foo.o foo.c`
 - ▶ preprocess and compile only, creating object file `foo.o`
- ▶ `cc -o foo.exec foo1.o foo2.o foo3.o`
 - ▶ link object files `foo1.o`, `foo2.o`, `foo3.o` and libraries to form executable `foo.exec`

Preprocessor directives

▶ `#include "filename" or #include <filename>`

▶ insert contents of filename here

▶ `#define X some text`

▶ let X = "some text"

▶ `#define X`

▶ let X be the empty string — but still defined

▶ `#ifdef X`

if X is defined, include ...

...

`#endif`

▶ `#ifdef X`

if X is defined, include ...

...

`#else`

else include ...

...

`#endif`

▶ `#if defined(X) && Y>2`

Preprocessor example motivation: constants

The length in a C array declaration in file scope must be a **constant expression** ...

- ▶ `int a[100];` : good
- ▶ `const int n=100;`
`int a[n];` : may or may not work
- ▶ 100 is definitely a constant expression
- ▶ is `n` a constant expression? that's up to the C implementation
- ▶ your goal should be to write **portable code**
 - ▶ will work for any conforming C compiler
 - ▶ compile with `-pedantic` to see if you rely on any non-portable features

```
basie:tmp siegel$ cc -pedantic -std=c11 tmp.c
tmp.c:2:5: warning: size of static array must be an integer constant
        expression [-Wpedantic]
int a[n];
    ^
1 warning generated.
```

Preprocessor example usage: constants

This always works:

```
#define N 100
int a[N];
int main() {
    for (int i=0; i<N; i++)
        a[i] = i;
}
```

After preprocessing, the code above becomes

```
int a[100];
int main() {
    for (int i=0; i<100; i++)
        a[i] = i;
}
```


Defining preprocessor macros on the command line

Compiling with the flag...

- ▶ `-DX`

- ▶ equivalent to inserting `#define X` at the beginning of the file

- ▶ `-DX=blah`

- ▶ equivalent to inserting `#define X blah` at the beginning of the file

Example:

```
// the preprocessor macro N must be defined (length of array a)
int a[N];
int main() {
    for (int i=0; i<N; i++)
        a[i] = i;
}
```

Compile:

```
cc -pedantic -DN=100 tmp.c
```

Example: boolean flags controlling compilation

```
// the preprocessor macro N must be defined (length of array a)
// define DEBUG to see debugging output
#include <stdio.h>
int a[N];
int main() {
#ifdef DEBUG
    printf("Entering for loop with N=%d\n", N);
    fflush(stdout);
#endif
    for (int i=0; i<N; i++)
        a[i] = i;
#ifdef DEBUG
    printf("Exiting for loop.\n");
    fflush(stdout);
#endif
}
```

To compile a “debugging version” of this program:

```
cc -pedantic -DDEBUG -DN=100 tmp.c
```

Preprocessor: Function-like macros

- ▶ the macros above are called **object-like macros**
- ▶ you can also **#define function-like macros**

```
#define MAX(x,y) ((x)>=(y) ? (x) : (y))  
int main() {  
    int m = MAX(N, 10);  
}
```

expands to

```
int main() {  
    int m = ((N)>=(10) ? (N) : (10));  
}
```

- ▶ why the abundance of parentheses?

Function-like macros: beware the pitfalls!

```
#define ADD(x,y) x+y  
#define MUL(x,y) x*y
```

What do these expand to?

- ▶ `ADD(1,2)*3` $1+2*3 = 7$
- ▶ `MUL(2,3+4)` $2*3+4 = 10$

Better:

```
#define ADD(x,y) ((x)+(y))  
#define MUL(x,y) ((x)*(y))
```

Structure of a C program

After preprocessing, the program consists of a sequence of

- ▶ **declarations**
 - ▶ variables
 - ▶ `int x;`
 - ▶ types
 - ▶ `typedef double D;`
 - ▶ enumerations
 - ▶ `enum Color { RED=0, GREEN=1, BLUE=2 };`
 - ▶ function **prototypes**
 - ▶ `int sgn(double x);`
- ▶ **function definitions**
 - ▶ `int sgn(double x) {
 if (x>0)
 return 1;
 else if (x<0)
 return -1;
 else return 0;
}`

Types

- ▶ C is statically typed
- ▶ every variable and expression has a type that is known at compile time
 - ▶ “statically”
 - ▶ before running the program
- ▶ you should be able to read a program and identify the type of any expression
- ▶ a type can be **complete** or **incomplete**
- ▶ every complete type has a **size**
 - ▶ the number of bytes required to store one element of that type
- ▶ `sizeof(T)` is an expression that returns the size of type *T*
 - ▶ this is a positive integer
- ▶ examples
 - ▶ `sizeof(int)`
 - ▶ often 4, sometimes 8
 - ▶ must be big enough so that `int` can hold at least $-32767 \dots 32767$
 - ▶ `sizeof(float[10])`
 - ▶ size of an array of 10 `floats`
 - ▶ equals $10 * \text{sizeof(float)}$

char: the smallest type

- ▶ `char` : size is always one byte
 - ▶ a byte has at least 8 bits
 - ▶ the smallest addressable unit of memory
- ▶ `signed char` : includes at least $-128..127$
 - ▶ `signed` means the type includes positive and negative integers (and 0)
- ▶ `unsigned char` : includes at least $0..255$
 - ▶ `unsigned` means the type includes only nonnegative integers
- ▶ `char` is either `signed char` or `unsigned char`
 - ▶ which one is unspecified

Other integer types

- ▶ `(short | long) (signed | unsigned) int`
 - ▶ $3 * 2 = 6$ combinations, each a different type
 - ▶ `short signed int`, `short unsigned int`, `signed int`, ...
- ▶ abbreviations: `int` is optional, `signed` is the default
 - ▶ `short = short int = short signed int`
 - ▶ `long = long int = long signed int`
 - ▶ `int = signed int`
 - ▶ `unsigned short = unsigned short int`
 - ▶ `unsigned = unsigned int`
 - ▶ `unsigned long = unsigned long int`
- ▶ sizes
 - ▶ the C Standard specifies **minimum** ranges for each of these types
 - ▶ also `short` \leq "medium" \leq `long`
- ▶ `_Bool`
 - ▶ consists of exactly 0 and 1
 - ▶ a subtype of `int`

Floating types

- ▶ `float`

- ▶ floating point
- ▶ typically, 4 bytes = 32 bits

- ▶ `double`

- ▶ double precision floating point
- ▶ at least as precise as `float`
- ▶ typically, 8 bytes = 64 bits

Simple declarations

For these basic types

- ▶ syntax: type-name variable-name ;
- ▶ can declare multiple variables of the same type
- ▶ an initializer is optional

Examples:

- ▶ `int x;`
- ▶ `double y;`
- ▶ `unsigned long z;`
- ▶ `int x, y;`
- ▶ `int x = 3;`
- ▶ `int x=3, y, z=-17;`

Array types: declaration

Declaration

- ▶ if $T(x)$ declares x to have type T
- ▶ then $T(a[])$ declares a to have type *array-of- T*
- ▶ and $T(a[n])$ declares a to have type *array-of-length- n -of- T*

Declaration examples

- ▶ `double a[]`
 - ▶ $T(x) = \text{double } x$
 - ▶ declares x to have type `double`
 - ▶ $T(a[]) = \text{double } a[]$
 - ▶ declares a to have type *array-of-double*
 - ▶ incomplete array type
- ▶ `double a[n]`
 - ▶ $T(x) = \text{double } x$
 - ▶ $T(a[n]) = \text{double } a[n]$
 - ▶ declares a to have type *array-of-length- n -of-double*
 - ▶ complete array type

Complete and incomplete types

- ▶ in some places, a complete array type is required
 - ▶ whenever space must be allocated for the array
 - ▶ an ordinary (not parameter) declaration of a local or global variable of array type
- ▶ in other places a complete or incomplete array type may be used
 - ▶ a parameter declaration of array type
 - ▶ as the base type in a pointer type
- ▶ the element type of an array type must be a complete type
 - ▶ in multi-dimensional arrays, only the first length can be unspecified
 - ▶ `int a[] [10] [20];`
 - ▶ incomplete type, but complete element type
 - ▶ OK sometimes (e.g., function parameter)
 - ▶ `int a[10] [];`
 - ▶ incomplete element type – bad

Array example: simple 2d-array

```
#define N 2
#define M 3
int a[N][M];
int main() {
    // initialize...
    for (int i=0; i<N; i++)
        for (int j=0; j<M; j++)
            a[i][j] = i*M+j;
    // print...
    for (int i=0; i<N; i++) {
        for (int j=0; j<M; j++)
            printf("%d ", a[i][j]);
        printf("\n");
    }
}
```

Compile and execute:

```
basie:tmp siegel$ cc tmp.c
basie:tmp siegel$ ./a.out
0 1 2
3 4 5
basie:tmp
```

Exercise: transpose a square matrix

```
#define N 3
int a[N][N];
int main() {
    // initialize...
    for (int i=0; i<N; i++)
        for (int j=0; j<N; j++)
            a[i][j] = i*N+j;

    // INSERT: in-place transpose

    // print...
    for (int i=0; i<N; i++) {
        for (int j=0; j<N; j++)
            printf("%d ", a[i][j]);
        printf("\n");
    }
}
```

```
basie:tmp siegel$ ./a.out
0 3 6
1 4 7
2 5 8
```

Pointers

- ▶ a pointer is the address of a memory location
- ▶ pointers are **first-class objects** in C
- ▶ there are pointer types
- ▶ a pointer can be assigned using =
- ▶ a pointer can be passed as an argument in a function call
- ▶ a pointer can be returned by a function
- ▶ there are operations which consume pointers and return pointers
- ▶ a pointer is just like any other kind of data

Pointer types

- ▶ declaration
 - ▶ if $T(x)$ declares x to have type T
 - ▶ then $T(*p)$ declares p to have type *pointer-to- T*
- ▶ declaration examples
 - ▶ `double *p`
 - ▶ $T(x) = \text{double } x$
 - ▶ $T(*p) = \text{double } *p$
 - ▶ p has type *pointer-to-double*
 - ▶ `unsigned long int *p`
 - ▶ $T(x) = \text{unsigned long int } x$
 - ▶ $T(*p) = \text{unsigned long int } *p$
 - ▶ p has type *pointer-to-unsigned-long-int*

Pointer operations

There are two basic operations on pointers:

- ▶ **address-of** (`&`)

- ▶ given a variable, returns the address of that variable
- ▶ if `x` has type T then `&x` has type pointer-to- T
- ▶ example
 - ▶ `int x;`
`int *p = &x; // address of x`

- ▶ **dereference** (`*`)

- ▶ given a pointer, returns the value stored at that address
- ▶ if `p` has type pointer-to- T then `*p` has type T
- ▶ example
 - ▶ `int x = 5;`
`int *p = &x;`
`int y = 2 * (*p); // 10`

Pointer operations, cont


- ▶ `*p` can also be used on the left-hand side of an assignment

```
double x = 3.1415;  
double *p = &x;  
*p = 2.71828;  
printf("%lf", x); // 2.71828
```

Pointers into arrays

- ▶ you can also take the address of array elements

```
float a[10];  
float *p = &a[5];  
*p = 17;
```



Pointer into 2d-array

```
int a[2][3];  
int *p = &a[0][2];  
*p = 13;
```

