

# CISC 372: Parallel Computing

## Performance

Stephen F. Siegel


Department of Computer and Information Sciences  
University of Delaware

## Performance: definition

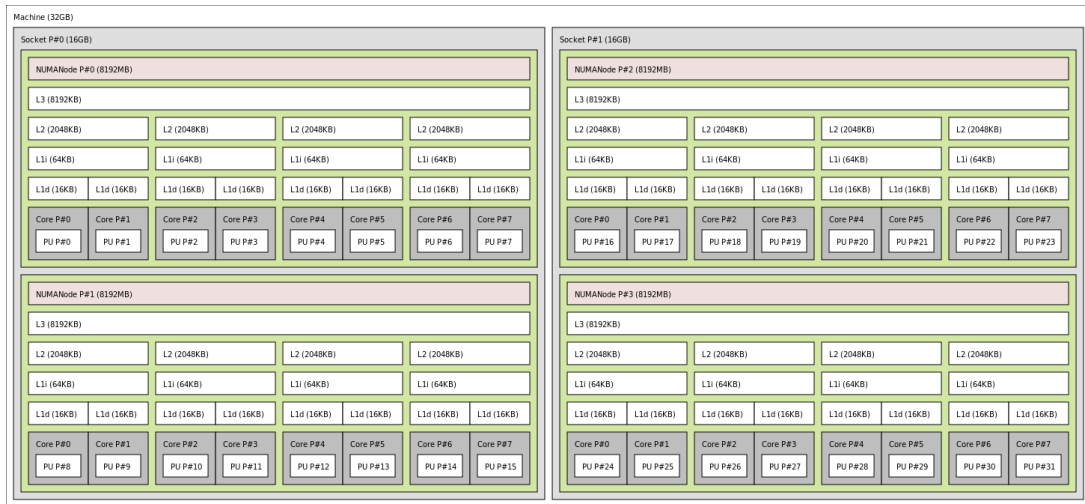
- ▶ how efficiently resources are used to solve a problem
- ▶ resources?
  - ▶ memory
  - ▶ energy
  - ▶ time

# Factors that affect performance: effective use of memory hierarchy

- ▶ modern CPUs have a hierarchy of **data caches** between CPU and memory
  - ▶ L1 cache: closest to core, very fast connection to registers (typical size: 32 KB/core)
  - ▶ L2 cache: further than L1, bigger, slower (256 KB/core)
  - ▶ L3 cache: further than L2, bigger, slower (2 MB/core)
  - ▶ DRAM: very slow



# Memory hierarchy: AMD Bulldozer server



[https://en.wikipedia.org/wiki/CPU\\_cache](https://en.wikipedia.org/wiki/CPU_cache)

## Memory hierarchy: example: matrix-vector multiplication

$$\begin{bmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \\ a_{30} & a_{31} & a_{32} \end{bmatrix} \times \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} a_{00}x_0 + a_{01}x_1 + a_{02}x_2 \\ a_{10}x_0 + a_{11}x_1 + a_{12}x_2 \\ a_{20}x_0 + a_{21}x_1 + a_{22}x_2 \\ a_{30}x_0 + a_{31}x_1 + a_{32}x_2 \end{bmatrix}$$

Layout of  $a$  in memory:

$a_{00}$	$a_{01}$	$a_{02}$	$a_{10}$	$a_{11}$	$a_{12}$	$a_{20}$	$a_{21}$	$a_{22}$	$a_{30}$	$a_{31}$	$a_{32}$
----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------	----------

- ▶ see `colmaj.c`:  $a$  is  $N \times N$  array of doubles,  $N = 20,000$ 
  - ▶ consider accesses to  $a$
  - ▶  $a[0][0]$ ,  $a[1][0]$ ,  $a[2][0]$ , ...
  - ▶ these are separated by  $20,000 * \text{sizeof}(\text{double})$  bytes!
  - ▶ each access loads into cache an entire block (cache line) containing the requested location
- ▶ see `rowmaj.c`
  - ▶ functionally equivalent to `colmaj.c`

## Factors that affect performance: compiler optimizations

Compilers can transform programs in myriad ways to use resources more effectively. . .

- ▶ function inlining; loop fission, loop fusion; loop interchange; loop unrolling; common subexpression elimination; constant folding, propagation . . .

Tradeoffs: more optimization generally entails. . .

- ▶ longer compile time
- ▶ larger generated code size
- ▶ program gets harder to debug
- ▶ greater sensitivity to undefined behavior (but you shouldn't use any undefined behavior!)
- ▶ generated code might actually get slower

Most compilers present a few pre-packaged optimization levels:

- ▶ `-O0`: little optimization, the default; `-Og`: recommended for debugging
- ▶ `-O1`, `-O2`, `-O3`: increasingly more optimizations applied
- ▶ see <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>

# Measuring performance of a parallel program

- ▶ to measure performance of a parallel program, you need a **baseline**
- ▶ baseline: a “similar” sequential program
  - ▶ using same inputs and other parameters to the extent possible
- ▶ different notions of **similar** are possible
  - ▶ do you choose the best possible sequential algorithm that solves the problem?
  - ▶ or the parallel program with  $-n$  1 (one process)?
  - ▶ these are often very different!
  - ▶ you must always specify the baseline

# Speedup

Let

$T_{\text{seq}}$  = time to run sequential baseline

$T_{\text{par}}$  = time to run parallel program

Then

$$\text{Speedup} = \frac{T_{\text{seq}}}{T_{\text{par}}}$$

- ▶ higher speedup is better
- ▶ if seq took 10 seconds and par took 2 seconds, speedup is 5
  - ▶ “parallel program is 5x faster than sequential”
  - ▶ with those particular inputs and nprocs



## Speedup as a function of `nprocs`

- ▶ hopefully: speedup will change (increase!) with `nprocs`
- ▶ **ideal case**: speedup = `nprocs`
  - ▶ double the number of procs, cut the execution time in half
- ▶ **reality**: rarely that good
  - ▶ communication time (sending messages)
  - ▶ synchronization time (procs have to sit around waiting, e.g., at a Barrier)
  - ▶ redundant work (two procs compute the same thing)
- ▶ after some point adding more processes no longer increases speedup
  - ▶ sort a list of  $10^{12}$  elements
    - ▶ doubtful you can improve speedup when `nprocs`  $> 10^{12}$
  - ▶ this is always the case for fixed problem size
- ▶ in best case, speedup may be approximately linear over some range of `nprocs`, but never as `nprocs`  $\rightarrow \infty$

# Amdahl's Law

- ▶ typically some part of the code cannot be parallelized
- ▶ “inherently sequential”
- ▶ example: `diffuse1d` writes data to the screen
  - ▶ there is only one screen: no way to parallelize that part
- ▶ example
  - ▶ say in sequential program, 10% of time is spent doing “inherently sequential” work
  - ▶ the other 90% can be parallelized
  - ▶ even if a parallel program were PERFECT with unlimited resources, the best it could do is reduce the 90% to 0.
  - ▶ the 10% time would be unchanged
  - ▶ therefore the best possible parallel time is  $(1/10) * T_{\text{seq}}$
  - ▶ best possible speedup is 10 :-(
- ▶ in general, if inherently sequential fraction of original program is  $r$ 
  - ▶ then speedup  $< 1/r$

## Weak Scaling vs. Strong Scaling

- ▶ Amdahl assumes inputs are held constant as **nprocs** increases
  - ▶ “strong scaling”
- ▶ Amdahl shows that strong scaling is problematic
  - ▶ there is always some point after which adding more procs cannot help
- ▶ **does this really reflect how people use parallel programs?**
- ▶ Gustafson (1988) said **no**
  - ▶ most users do **not** have some fixed problem size and ask *how fast can I make it?*
  - ▶ instead, the more processors you give them, the **bigger** they will make the problem size
  - ▶ almost every problem in science and engineering benefits from increased resolution or scale
- ▶ so a more useful measure of performance **increases problem size** with **nprocs**
  - ▶ “weak scaling”
- ▶ parallelization is more effective with weak scaling than with strong scaling

# Weak vs. Strong Scaling

- ▶ strong scaling: baseline is constant as nprocs increases
  - ▶ always comparing against sequential run on fixed problem size
  - ▶ speedup is bounded
- ▶ strong scaling examples
  - ▶ fix list of length  $10^6$ ; compare sequential time to sort vs. parallel time to sort with  $p$  procs
  - ▶ fix  $nx = 10^3$ ; compare sequential diffusion1d vs. parallel diffusion1d with  $p$  procs
    - ▶ note  $nxl$ , the amount of data per process, decreases as  $p$  increases
- ▶ weak scaling: baseline increases with nprocs
  - ▶ problem size of sequential program increases with nprocs
  - ▶ it is possible for speedup  $\rightarrow \infty$  as nprocs  $\rightarrow \infty$
- ▶ weak scaling examples
  - ▶ for  $p > 0$ , compare sequential time to sort list of length  $10^6 p$  with parallel time using  $p$  procs
  - ▶ for  $p > 0$ , compare sequential diffusion1d with  $nx = 10^3 p$  vs. parallel diffusion1d with  $p$  procs
    - ▶ note  $nxl = 10^3$  is held constant as  $p$  increases

# Efficiency

$$\text{efficiency} = \frac{\text{speedup}}{\text{nprocs}} = \frac{T_{\text{seq}}}{T_{\text{par}} * \text{nprocs}}$$

- ▶ efficiency is “speedup per process”
- ▶ Amdahl says that for strong scaling:
  - ▶  $\text{efficiency} \rightarrow 0$  as  $\text{nprocs} \rightarrow \infty$
- ▶ for weak scaling, in best case it is possible:
  - ▶  $\text{efficiency} \rightarrow 1$  as  $\text{nprocs} \rightarrow \infty$
- ▶ more common: something between 0 and 1

# Automating performance experiments

- ▶ see [exp/sat\\_strong](#) in public course repo
- ▶ a strong scaling experiment of MPI SAT solver
- ▶ [sat\\_mpi.c](#) has been altered to
  - ▶ print to [stdout](#) only the number of processes and time
  - ▶ other things are sent to [stderr](#)

```
if (rank == 0) {  
    const double time = MPI_Wtime() - start_time;  
  
    fprintf(stderr, "Number of solutions = %u.  Time = %lf.\n",  
            nsolutions, time);  
    fflush(stderr);  
    printf("%d %lf\n", nprocs, time);  
}
```

## Data file generated by SAT performance experiment

- ▶ the `Makefile` executes `sat_mpi.exec` with 1, 2, 4, 8, 16, 32 procs
- ▶ the results are accumulated in a file `sat_mpi.dat`:

```
1 42.693483
2 29.942159
4 16.342128
8 9.844605
16 5.327622
32 2.452447
```

nprocs	time	speedup	efficiency
1	43.75	1.00	1.00
2	30.46	1.44	0.72
4	16.47	2.66	0.66
8	8.68	5.04	0.63
16	4.89	8.95	0.56
32	2.52	17.36	0.54

## Graphing data with `gnuplot`


- ▶ free, open-source command-line tool for creating graphs
- ▶ <http://www.gnuplot.info>
- ▶ command: `gnuplot sat_mpi.gnu`

```
set terminal pdf
set output "sat_mpi.pdf"
set xlabel center "Number of processes"
set ylabel center "time (seconds)"
set xr [0:32]
set yr [0:45]
plot "sat_mpi.dat" using 1:2 title 'MPI' with linespoints
```

- ▶ meaning of `using 1:2`
  - ▶ use column 1 of the data file for the x-coordinates
  - ▶ use column 2 of the data file for the y-coordinates



## PDF file resulting from SAT scaling experiment



# Makefile for SAT performance experiment

```
ROOT = ../../
include $(ROOT)/common.mk
NAME = sat_mpi
all: $(NAME).exec
$(NAME).exec: $(NAME).c Makefile
    $(MPICCC) -o $@ $<
$(NAME).dat: $(NAME).exec
    $(MPIRUN) -n 1 ./$(NAME).exec > $(NAME).dat
    $(MPIRUN) -n 2 ./$(NAME).exec >> $(NAME).dat
    $(MPIRUN) -n 4 ./$(NAME).exec >> $(NAME).dat
    $(MPIRUN) -n 8 ./$(NAME).exec >> $(NAME).dat
    $(MPIRUN) -n 16 ./$(NAME).exec >> $(NAME).dat
    $(MPIRUN) -n 32 ./$(NAME).exec >> $(NAME).dat
graphs:
    gnuplot $(NAME).gnu
.PHONY: all graphs
```

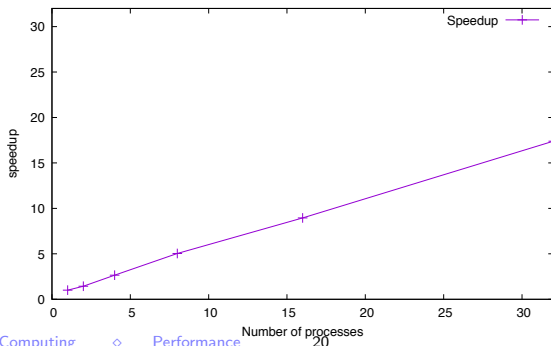
# Using gnuplot

Much more is possible. . .

- ▶ graph speedup
- ▶ graph efficiency
- ▶ graph multiple plots in one picture
  - ▶ e.g.: sequential vs. MPI vs. OpenMP

# Graphing speedup with gnuplot

```
set output "sat_speedup.pdf"
set xlabel center "Number of processes"
set ylabel center "speedup"
set xr [0:32]
set yr [0:32]
first(x) = ($0 > 0 ? base : base = x)
plot "sat_mpi.dat" using 1:(first($2), base/$2) title 'Speedup' with linespoints
```



# Graphing efficiency with gnuplot

```
set output "sat_efficiency.pdf"
set xlabel center "Number of processes"
set ylabel center "efficiency"
set xr [0:32]
set yr [0:1]
first(x) = ($0 > 0 ? base : base = x)
plot "sat_mpi.dat" using 1:(first($2), base/($2*$1)) title 'Efficiency' with linespoints
```

