

CISC 372: Parallel Programming

MPI Point-to-Point Operations


Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

Point to Point Operations


- ▶ for sending a message from one process to another process
- ▶ sending process issues a **send** instruction
- ▶ receiving process issues a **receive** instruction
- ▶ can be considered “lower-level” than collective operations
- ▶ all collective operations can be implemented using point-to-points
 - ▶ but quality MPI implementations will provide better performance for collectives
- ▶ “push” model (like the mail)
 - ▶ sending process specifies destination
 - ▶ receiving process may or may not specify source

Message channels: conceptual framework



- ▶ the state of a communicator with 3 procs
- ▶ every communicator is isolated — has its own state
 - ▶ messages from one communicator are never picked up by an operation from a different communicator
- ▶ between any 2 procs, there is a **p2p message channel**
 - ▶ including from proc to itself (rarely used)
- ▶ **send** enqueues message
- ▶ **recv** dequeues message
- ▶ mostly a FIFO queue

Tags



- ▶ each message has a **tag**
- ▶ an **int** specified by the sender
- ▶ the receiver **may** specify a tag
 - ▶ or can specify “any tag”
- ▶ if P2 issues recv from P0 with tag 2
 - ▶ P2 will receive message 1
- ▶ if P2 issues recv from P0 with tag 1
 - ▶ P2 will receive message 2
 - ▶ the first (oldest) message in queue with matching tag
- ▶ if P2 issues recv from P0 with “any tag”
 - ▶ P2 will receive message 1

MPI_Send

`MPI_Send(buf, count, datatype, dest, tag, comm)`

`buf` address of send buffer (`void*`)
`count` number of elements in buffer (`int`)
`datatype` data type of elements in buffer (`MPI_Datatype`)
`dest` rank of destination process (`int`)
`tag` integer to attach to message `envelope` (`int`)
`comm` communicator (`MPI_Comm`)

- ▶ message `envelope`
 - ▶ source rank
 - ▶ destination rank
 - ▶ tag
 - ▶ communicator
- ▶ tag can be used by receiver to select which message to receive

MPI_Recv

MPI_Recv(buf, count, datatype, source, tag, comm, status)

buf	address of receive buffer (<code>void*</code>)
count	number of elements in buffer (<code>int</code>)
datatype	data type of elements in buffer (<code>MPI_Datatype</code>)
source	rank of source process (<code>int</code>)
tag	tag of message to receive (<code>int</code>)
comm	communicator (<code>MPI_Comm</code>)
status	pointer to status object (<code>MPI_Status*</code>)

- ▶ `count` must be at least as large as count of incoming message
 - ▶ otherwise, **undefined behavior**
- ▶ `status`: object to store envelope information on received message
 - ▶ source, tag, count
 - ▶ if you don't need it, use `MPI_STATUS_IGNORE`
- ▶ why would you need to know `source` and `tag` when you already specified them?

Example: p2p.c

```
#include<stdio.h>
#include<mpi.h>
int main() {
    int message, rank;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        message = 173;
        MPI_Send(&message, 1, MPI_INT, 1, 9, MPI_COMM_WORLD);
    } else if (rank == 1) {
        MPI_Recv(&message, 1, MPI_INT, 0, 9, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Proc 1 received: %d\n", message);
    }
    MPI_Finalize();
}
```

```
> mpiexec -n 4 ./p2p.exec
```

```
Proc 1 received: 173
```

Example: using different tags: tags.c

```
/* tags.c: demonstration of receiving messages out of order using tags. Note that
   this program is not safe --- technically, it could deadlock. But if it does not
   deadlock, the messages will be received in the reverse order. */
#include<stdio.h>
#include<mpi.h>
int main() {
    int message, rank;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        message = 1; MPI_Send(&message, 1, MPI_INT, 1, 1, MPI_COMM_WORLD); // tag=1
        message = 2; MPI_Send(&message, 1, MPI_INT, 1, 2, MPI_COMM_WORLD); // tag=2
    } else if (rank == 1) {
        MPI_Recv(&message, 1, MPI_INT, 0, 2, MPI_COMM_WORLD, MPI_STATUS_IGNORE); // tag=2
        printf("Proc 1 received: %d\n", message);
        MPI_Recv(&message, 1, MPI_INT, 0, 1, MPI_COMM_WORLD, MPI_STATUS_IGNORE); // tag=1
        printf("Proc 1 received: %d\n", message);
    }
    MPI_Finalize();
}
```


MPI_ANY_TAG

- ▶ a recv can use `MPI_ANY_TAG` for the tag argument
- ▶ receive a message from sender with “any tag”
- ▶ it will always match the **oldest** message from the sender
- ▶ execution is **deterministic** — one and only one thing can happen

Example: using MPI_ANY_TAG: anytag.c

```
/* anytag: the messages will be received in the order sent.  The MPI_ANY_TAG recv
   must match the oldest message sent from proc 0 */
#include<stdio.h>
#include<mpi.h>
int main() {
    int message, rank;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        message = 1;
        MPI_Send(&message, 1, MPI_INT, 1, 1, MPI_COMM_WORLD); // tag=1
        message = 2;
        MPI_Send(&message, 1, MPI_INT, 1, 2, MPI_COMM_WORLD); // tag=2
    } else if (rank == 1) {
        MPI_Recv(&message, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Proc 1 received: %d\n", message);
        MPI_Recv(&message, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Proc 1 received: %d\n", message);
    }
    MPI_Finalize();
}
```

Getting the status

`status` is a C `struct`

- ▶ getting the rank of the source
 - ▶ `status.MPI_SOURCE`
- ▶ getting the tag of the message
 - ▶ `status.MPI_TAG`
- ▶ getting the error code
 - ▶ `status.MPI_ERROR`
- ▶ getting the size (“count”) of the message
 - ▶ not simply a field in the struct
 - ▶ need to use function `MPI_Get_count`

Example: status.c

```
#include<string.h>
#include<stdio.h>
#include<mpi.h>

int main() {
    char message[100];
    int rank;
    MPI_Status status;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        strcpy(message,"Hello, from proc 0!");
        MPI_Send(message, strlen(message)+1, MPI_CHAR, 1, 99, MPI_COMM_WORLD);
    } else if (rank == 1) {
        MPI_Recv(message, 100, MPI_CHAR, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
        printf("Proc 1 received: \"%s\"\n", message);
        printf("source=%d tag=%d \n", status.MPI_SOURCE, status.MPI_TAG);
    }
    MPI_Finalize();
}
```

status.c output

Note that in C, a string is a sequence of `char` ending with the “null terminating char” `'\0'`. The number of characters in the string is therefore `strlen(message) + 1 = 19 + 1 = 20`.

```
> mpiexec status.exec  
Proc 1 received: "Hello, from proc 0!"  
source=0 tag=99
```

MPI_Get_count

`MPI_Get_count(status, datatype, count)`

`status` pointer to status object (`MPI_Status*`)
`datatype` data type of elements received (`MPI_Datatype`)
`count` pointer to variable in which to return result (`int*`)

- ▶ should only be called after `status` has been filled in by receive
- ▶ `datatype` should be same as used in receive
- ▶ sets `count` to the number of elements received
- ▶ **note**
 - ▶ `count` specified in receive statement and message `count` can differ
 - ▶ receive buffer must be big enough to hold incoming message
 - ▶ memory in receive buffer after message count will not be altered

Example: getting the count: `count.c`

The following lines are added to proc 1:

```
int count;
MPI_Get_count(&status, MPI_CHAR, &count);
printf("source=%d tag=%d count=%d\n",
       status.MPI_SOURCE, status.MPI_TAG, count);
```

This sets `count` to the actual number of characters (`MPI_CHAR`) received.

```
> mpiexec -n 4 ./count.exec
Proc 1 received: "Hello, from proc 0!"
source=0 tag=99 count=20
```

Note the null terminating character is counted.

Synchronization and deadlock

- ▶ a receive operation must **block** until a matching message arrives
- ▶ this can lead to **deadlocks** if you are not careful; see [deadlock.c](#)

```
#include<stdio.h>
#include<mpi.h>
int main() {
    int message, rank;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        message = 173;
        printf("Proc 0: was I supposed to do something?\n");
    } else if (rank == 1) {
        MPI_Recv(&message, 1, MPI_INT, 0, 9, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Proc 1 received: %d\n", message);
    }
    MPI_Finalize();
}
```

```
mpiexec -n 4 ./deadlock.exec
Proc 0: was I supposed to do something?
^C[mpiexec@basie.local] Sending Ctrl-C to processes as requested
```


Synchronization and potential deadlock

- ▶ a send operation ...
 - ▶ may complete even if a matching receive operation has not been executed
 - ▶ the message will be stored in a system buffer (channel)
 - ▶ or it may block until a matching receive is available
 - ▶ the message can then be copied directly from send buffer to recv buffer
- ▶ the choice is up to the MPI implementation
- ▶ the decision can be made differently at each send operation
- ▶ you cannot assume anything
- ▶ a correct program will behave correctly regardless of how this decision is made

Example `may_deadlock.c`: a potential deadlock

```
#include<stdio.h>
#include<mpi.h>

int main() {
    int message, rank;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        message = 173;
        MPI_Send(&message, 1, MPI_INT, 1, 9, MPI_COMM_WORLD);
    } else if (rank == 1) {
        printf("Proc 1: was I supposed to do something?\n");
    }
    MPI_Finalize();
}
```

Exchanging data

- ▶ suppose two processes wish to **exchange some data**
 - ▶ proc 0 wants to send something to proc 1, and
 - ▶ proc 1 wants to send something to proc 0
- ▶ very common scenario
- ▶ how to it safely?
 - ▶ must be correct
 - ▶ must not deadlock

Exchange 1: Incorrect: will deadlock!

- ▶ both procs try to receive before sending

```
int main() {
    int rank, myNumber, otherNumber;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        myNumber = 10;
        MPI_Recv(&otherNumber, 1, MPI_INT, 1, 9, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        MPI_Send(&myNumber, 1, MPI_INT, 1, 9, MPI_COMM_WORLD);
    } else if (rank == 1) {
        myNumber = 20;
        MPI_Recv(&otherNumber, 1, MPI_INT, 0, 9, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        MPI_Send(&myNumber, 1, MPI_INT, 0, 9, MPI_COMM_WORLD);
    }
    printf("Process %d: received %d\n", rank, otherNumber);
    MPI_Finalize();
}
```

Exchange 2: Unsafe: may deadlock!

- ▶ both procs send before receiving — what if MPI tries to execute both sends synchronously?

```
int main() {
    int rank, myNumber, otherNumber;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        myNumber = 10;
        MPI_Send(&myNumber, 1, MPI_INT, 1, 99, MPI_COMM_WORLD);
        MPI_Recv(&otherNumber, 1, MPI_INT, 1, 99, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    } else if (rank == 1) {
        myNumber = 20;
        MPI_Send(&myNumber, 1, MPI_INT, 0, 99, MPI_COMM_WORLD);
        MPI_Recv(&otherNumber, 1, MPI_INT, 0, 99, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    }
    printf("Process %d: received %d\n", rank, otherNumber);
    MPI_Finalize();
}
```

Exchange 3: Correct: procs alternate

- ▶ one proc sends, then receives; the other proc receives, then sends

```
int main() {
    int rank, myNumber, otherNumber;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        myNumber = 10;
        MPI_Send(&myNumber, 1, MPI_INT, 1, 99, MPI_COMM_WORLD);
        MPI_Recv(&otherNumber, 1, MPI_INT, 1, 99, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    } else if (rank == 1) {
        myNumber = 20;
        MPI_Recv(&otherNumber, 1, MPI_INT, 0, 99, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        MPI_Send(&myNumber, 1, MPI_INT, 0, 99, MPI_COMM_WORLD);
    }
    printf("Process %d: received %d\n", rank, otherNumber);
    MPI_Finalize();
}
```

Exchanging with `MPI_Sendrecv`

- ▶ this situation is so common, MPI provides a function to deal with it
- ▶ `MPI_Sendrecv` combines one send and one receive operation into a single command
- ▶ both operations execute concurrently

MPI_Sendrecv

```
MPI_Sendrecv(sbuf, scount, stype, dest, stag,  
             rbuf, rcount, rtype, source, rtag,  
             comm, status)
```

sbuf	address of send buffer (<code>void*</code>)
scount	number of elements in send buffer (<code>int</code>)
stype	data type of elements in sbuf (<code>MPI_Datatype</code>)
dest	rank of destination process (<code>int</code>)
stag	integer to attach to message <code>envelope</code> (<code>int</code>)
rbuf	address of receive buffer (<code>void*</code>)
rcount	length of receive buffer (<code>int</code>)
rtype	data type of elements to be received (<code>MPI_Datatype</code>)
source	rank of sending process (<code>int</code>)
rtag	tag of message to receive (<code>int</code>)
comm	communicator (<code>MPI_Comm</code>)
status	pointer to status object for receive (<code>MPI_Status*</code>)

Semantics and uses of MPI_Sendrecv

- ▶ combines a send statement and a receive statement into one statement
- ▶ both operations post simultaneously
- ▶ as if two threads are spawned, one to manage the send, the other the receive
- ▶ the operation completes only after both the send and receive complete
- ▶ solves the deadlocking problem for data exchange
- ▶ cyclic exchange
 - ▶ $0 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 0$
 - ▶ process of rank i
 - ▶ sends to $i + 1$ (modulo numProcs)
 - ▶ receives from $i - 1$ (modulo numProcs)
- ▶ shift
 - ▶ $0 \rightarrow 1 \rightarrow 2 \rightarrow 3$
 - ▶ proc 0 only sends
 - ▶ proc `nprocs` - 1 only receives
 - ▶ or use `MPI_PROC_NULL`

Exchange 4: Correct: MPI_Sendrecv

```
int main() {
    int rank, myNumber, otherNumber;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        myNumber = 10;
        MPI_Sendrecv(&myNumber, 1, MPI_INT, 1, 99, &otherNumber, 1, MPI_INT, 1, 99,
                     MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    } else if (rank == 1) {
        myNumber = 20;
        MPI_Sendrecv(&myNumber, 1, MPI_INT, 0, 99, &otherNumber, 1, MPI_INT, 0, 99,
                     MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    }
    if (rank < 2) printf("Process %d: received %d\n", rank, otherNumber);
    MPI_Finalize();
}
```

Cyclic exchange: cycle.c

```
#include<stdio.h>
#include<mpi.h>

int main() {
    int nprocs, rank;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    const int right = (rank + 1)%nprocs, left = (rank + nprocs - 1)%nprocs;
    int rbuf, sbuf = 100 + rank;
    MPI_Sendrecv(&sbuf, 1, MPI_INT, right, 0, &rbuf, 1, MPI_INT, left, 0,
                 MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Proc %d: received %d\n", rank, rbuf);
    MPI_Finalize();
}
```

- note use of `rank + nprocs - 1` to avoid a negative argument to modulo operator

Shift exchange: shift.c

```
#include<stdio.h>
#include<mpi.h>

int main() {
    int nprocs, rank;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    const int right = rank < nprocs - 1 ? rank + 1 : MPI_PROC_NULL,
        left = rank > 0 ? rank - 1 : MPI_PROC_NULL;
    int rbuf, sbuf = 100 + rank;
    MPI_Sendrecv(&sbuf, 1, MPI_INT, right, 0, &rbuf, 1, MPI_INT, left, 0,
        MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    if (rank > 0) printf("Proc %d: received %d\n", rank, rbuf);
    MPI_Finalize();
}
```

► a send or receive to **MPI_PROC_NULL** is a no-op

Semantics: Non-interaction with collectives

- ▶ an MPI program can use both point-to-point and collective operations
- ▶ point-to-point and collective operations **exist in two separate universes**
 - ▶ there is no “matching” between p2p and collective operations
 - ▶ a message sent by a p2p can never be received by a collective
 - ▶ a message sent by a collective can never be received by a p2p