

CISC 372: Parallel Computing Threads

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

Operating System Concepts: Processes and Threads

- ▶ a thread is a “light-weight process”
- ▶ a process consists of ...
 - ▶ memory for the **heap**
 - ▶ **resources** the OS has allocated for the process: file descriptors, ...
 - ▶ environment variables
 - ▶ security information: what can this process access?
 - ▶ state information: is this process running, waiting, ...?
 - ▶ private address space (cannot be accessed directly by other processes)
 - ▶ In Unix, **sockets** can be used for inter-process communication
 - ▶ **multiple threads**
- ▶ a **thread** consists of ...
 - ▶ a stack (one frame for each function call)
 - ▶ program counter (reference to current position in the program)
- ▶ “thread” means “thread of control” in a program
- ▶ creating threads is “cheap” (compared to creating a whole process)
 - ▶ an OS can support many threads [see “Activity Monitor” in OS X]

Programming with threads in C

- ▶ originally, C did not support threads directly
- ▶ POSIX threads = **Pthreads**
 - ▶ an API for thread programming in C
 - ▶ POSIX is the “Unix” Standard (Portable Operating System Interface)
 - ▶ IEEE Std. 1003.1-2017, <https://pubs.opengroup.org/onlinepubs/9699919799/>
 - ▶ 4 volumes, thousands of pages
 - ▶ examples: Linux, BSD Unix (mostly conformant); macOS, AIX, HP-UX (certified)
 - ▶ by far the most common way to program threads in C
- ▶ **C11** = C Standard from 2011 (C18 is current version)
 - ▶ added support for threads
 - ▶ some new syntax, mostly a standard library
 - ▶ similar in many respects to Pthreads
 - ▶ should be portable to any conforming C implementation supporting threads (not just Unix)
 - ▶ very complicated “memory model” (in comparison with Pthreads)
 - ▶ compilers, text books, tutorial, ... still catching up
 - ▶ some criticism from the user community on technical aspects

Other programming languages

- ▶ C++11, C++14, ...: threads
 - ▶ the C11 thread model was basically ported from C++11
- ▶ Java
 - ▶ always had threads built into the language
 - ▶ Thread, synchronize, wait(), notify, ...
 - ▶ still not easy to use
- ▶ Ada
 - ▶ 1983, 1995, 2005, 2012
 - ▶ always had concurrency: tasks
 - ▶ very elegant concurrency model
- ▶ many functional languages
 - ▶ Racket
 - ▶ Haskell
 - ▶ ...

First, more C: Function types

- ▶ if $T(x)$ declares x to have type T
 - ▶ then $T(f())$ declares f to have type *function-returning- T*
- ▶ $T(f(t_1, t_2, \dots, t_n))$
 - ▶ where the t_i are type names
 - ▶ further declares f to consume inputs of the named types
 - ▶ parameter names may also be included in the declaration but have no meaning
- ▶ $T(f(\text{void}))$ indicates that f has 0 parameters
- ▶ $()$ binds tighter than $*$
 - ▶ `int *f();` : function-returning-pointer-to-int
 - ▶ `int (*f)();` : pointer-to-function-returning-int

Automatic conversion: functions

- ▶ an expression or parameter of type **function-returning-T**
 - ▶ is automatically converted to type **pointer-to-function-returning-T**
 - ▶ exceptions
 - ▶ operand of `sizeof`
 - ▶ operand of `&`
 - ▶ operand of `_Alignof`
- ▶ this is a ***function pointer***
 - ▶ in particular, the function used in a function call is really a function pointer
 - ▶ if you pass a function as an argument in a call
 - ▶ it is really a pointer to the function that you pass
- ▶ example
 - ▶ `int f(int);`
 - ▶ whenever `f` is used as an expression it is interpreted as a **pointer** to the function named `f`
 - ▶ `f(5)` : the `f` in this call is a pointer to the function

Exercise

Write a function `addmeup` which consumes

- ▶ a function `f` which consumes an `int` and returns an `int`

and returns

- ▶ $\sum_{i=0}^9 f(i)$

Pthreads

- ▶ program begins executing with a single thread of control — just like a sequential program
- ▶ new threads are created dynamically by calling `pthread_create`
- ▶ you specify the function to run in the new thread with a function pointer
- ▶ global variables are **shared**
 - ▶ all threads can read and write to them
 - ▶ this is how threads communicate (as opposed to message-passing)
 - ▶ this must be done very carefully!
- ▶ local variables are **private**
- ▶ various synchronization mechanisms are provided by Pthreads
 - ▶ wait until another thread terminates
 - ▶ block until a lock becomes available. . . release the lock
 - ▶ wait until some condition becomes *true*
- ▶ requires a compiler and OS supporting Pthreads
 - ▶ all standard Unix systems

pthread_t

- ▶ a **type**
- ▶ a “thread ID” or reference to a running thread
- ▶ an **opaque** object
 - ▶ you cannot perform any operations on something of type pthread_t
 - ▶ you can only use that thing as an argument to functions in the Pthreads library
 - ▶ you don't know how it is actually defined
 - ▶ it may be a pointer, or an int, or ...
 - ▶ completely up to the Pthread implementation

pthread_create

```
int pthread_create(pthread_t * thread,  
                  pthread_attr_t * attr,  
                  void *(*start_routine)(void*),  
                  void * arg);
```

pthread_create(thread, attr, start_routine, arg)

thread	where ID of new thread is returned (pthread_t*)
attr	thread attributes, can be NULL (pthread_attr_t*)
start_routine	function to run in new thread (void *(*)(void*))
arg	argument to provide to start_routine (void*)

Note: start_routine consumes a void* and returns a void*

pthread_join

```
int pthread_join(pthread_t thread, void **value_ptr);
```

```
pthread_join(thread, value_ptr)
```

<code>thread</code>	ID of thread to wait on (<code>pthread_t</code>)
<code>value_ptr</code>	address in which to store return val (<code>void**</code>)

- ▶ blocks until the specified thread terminates
- ▶ deallocates the thread resources
- ▶ similar to MPI_Wait

Hello, world!

See [hello/hello_pthread.c](#).

Notes:

- ▶ the start function must consume and return a `void*`
- ▶ any pointer type can be converted to a `void*` and back again
- ▶ result is guaranteed to be equal to original pointer (C Standard)
- ▶ in this example: an `int*` is used for each thread
- ▶ need a separate one for each thread so each can have their own ID
- ▶ some people will convert a `long` to a `void*` and back
 - ▶ `usually` works
 - ▶ but result is `not guaranteed` to be the original `long`
 - ▶ not safe or portable!