

CISC 372: Parallel Computing Threads, part 3: condition variables

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

Example: bank account

```
const int max = 10; // keep bal in 0..max
int bal = 0;
pthread_mutex_t mutex;
```

```
void * deposit_thread(void * arg) {
    while (1) {
        // WAIT UNTIL bal<10 ...
        pthread_mutex_lock(&mutex);
        bal++;
        pthread_mutex_unlock(&mutex);
    }
}
```

```
void * withdraw_thread(void * arg) {
    while (1) {
        // WAIT UNTIL bal>0 ...
        pthread_mutex_lock(&mutex);
        bal--;
        pthread_mutex_unlock(&mutex);
    }
}
```

- ▶ only want depositor to take the lock if `bal<10`
- ▶ only want withdrawer to take the lock if `bal>0`
- ▶ an example of the **producer-consumer** pattern

Bad solution

```
while (true) {  
    pthread_mutex_lock(&mutex);  
    if (bal > 0) break;  
    pthread_mutex_unlock(&mutex);  
}  
bal--;  
pthread_mutex_unlock(&mutex);
```

- ▶ functionally, this is correct
- ▶ performance-wise: disaster
 - ▶ thread is constantly spinning, rechecking `bal` repeatedly, unnecessarily
 - ▶ ...and taking and releasing lock
 - ▶ a thread that should be quietly waiting is instead constantly consuming resources (CPU)
 - ▶ if many threads do this: lock contention
 - ▶ performance grinds to a halt

Monitors

- ▶ the “**monitor**” is a standard solution to this problem
- ▶ a concurrency concept introduced by Per Brinch Hansen and C.A.R. Hoare in early 1970s
- ▶ used in many programming languages/APIs
- ▶ Concurrent Pascal (1974, Hansen)
- ▶ Java: `synchronize`, `wait()`, `notify()`, `notifyAll()`
- ▶ Pthreads: **condition variables** and **mutexes**
- ▶ monitor = condition variable + mutex

Condition variables

- ▶ a **condition variable** *c* is used with a mutex
- ▶ when a thread owns the mutex, it may want to wait until some condition holds (due to actions of other threads)
- ▶ it can do this by **waiting on *c***
- ▶ this relinquishes the locks and the thread goes to sleep
- ▶ other threads run
- ▶ at some point in future, another thread can issue a **notification on *c***
- ▶ the thread that is asleep may be notified
 - ▶ it wakes up and has the opportunity to regain the lock once the thread owning the lock relinquishes it
 - ▶ typically, after the thread wakes up, it will check some condition
 - ▶ if the condition holds, great, it continues running with the lock
 - ▶ otherwise, it waits again (loops are good for this)

Condition variables in Pthreads

- ▶ `pthread_cond_init(pthread_cond_t * cond, NULL)`
 - ▶ initialize a condition variable
- ▶ `int pthread_cond_destroy(pthread_cond_t * cond);`
 - ▶ destroy the previously initialized condition variable
- ▶ `int pthread_cond_signal(pthread_cond_t * cond);`
 - ▶ wake up one or more threads waiting on `cond`
- ▶ `int pthread_cond_broadcast(pthread_cond_t * cond);`
 - ▶ wake up all threads waiting on `cond`
- ▶ `int pthread_cond_wait(pthread_cond_t * cond, pthread_mutex_t * mutex);`
 1. release lock on `mutex`
 2. go to sleep
 3. when woken up: try to regain lock on `mutex`

Semantics of a condition variable `c`

- ▶ every thread is either **running**, **blocked** waiting for lock, or **asleep**
 - ▶ a sleeping thread is not contending for resources/consuming CPU cycles
 - ▶ note: I am using “asleep” here in a non-standard sense
- ▶ **state** of `c`: set of waiting threads (“**wait-set**”)
- ▶ **wait** involves 3 atomic operations:
 1. release lock on mutex, state changes from **running** to **asleep**, thread added to `c`’s wait-set
 2. when signaled: state changes from **asleep** to **blocked**, thread removed from `c`’s wait-set
 3. later the thread may regain the lock on **mutex**
 - ▶ just like any thread trying to unlock **mutex**
 - ▶ once lock has been obtained, the call to **wait** returns
- ▶ **signal**
 - ▶ changes state of one or more waiting threads as above, removes them from `c`’s wait-set
 - ▶ usually called from thread that owns lock on **mutex**, but not required by Pthreads
- ▶ **broadcast**: signals all waiting threads, `c`’s wait-set become empty

Typical pattern for using condition variables

```
obtain lock on mutex;
...
while (!expr) {
    wait on cond;
}
// at this point you know expr holds
// assuming expr can only be changed
// by a thread holding lock on mutex!
...
release lock on mutex;
```


Bank account: bank1.c

```
const int max = 10; // keep bal in 0..max
int bal = 0;
pthread_mutex_t mutex;
pthread_cond_t balLT10, balGT0;
```

```
void * deposit_thread(void * arg) {
    while (1) {
        pthread_mutex_lock(&mutex);
        while (!(bal<max))
            pthread_cond_wait(&balLT10, &mutex);
        // now I know bal<10 and I have the lock
        bal++;
        pthread_cond_signal(&balGT0);
        pthread_mutex_unlock(&mutex);
    }
}
```

```
void * withdraw_thread(void * arg) {
    while (1) {
        pthread_mutex_lock(&mutex);
        while (!(bal>0))
            pthread_cond_wait(&balGT0, &mutex);
        // now I know bal>0 and I have the lock
        bal--;
        pthread_cond_signal(&balLT10);
        pthread_mutex_unlock(&mutex);
    }
}
```

Generalized: `bank2.c`

- ▶ now allow multiple accounts, multiple depositors, multiple withdrawers
- ▶ a depositor randomly chooses an account and an amount
 - ▶ deposits the amount to the account (no waiting)
 - ▶ repeat forever
- ▶ a withdrawer randomly chooses an account and an amount
 - ▶ **waits for the balance to be at least the amount**
 - ▶ withdraws the amount from the account
 - ▶ repeat forever
- ▶ command line args: number of accounts, number of depositors, number of withdrawers
- ▶ solution
 - ▶ one mutex and one condition variable for each account
 - ▶ mutex guards all accesses to the account balance
 - ▶ condition variable signals whenever a deposit is made to the account
 - ▶ depositor signals every time it makes a deposit to the account
 - ▶ withdrawer waits, and upon being signaled, checks the balance

Application: concurrency flags

- ▶ a **flag** is a boolean variable
- ▶ a **concurrency flag** is a shared boolean variable used in a particular disciplined way
 - ▶ also known as a “binary semaphore”
- ▶ concurrency **flags** are basic concurrency building blocks
- ▶ can be used to construct all kinds of complex synchronization patterns and data structures
 - ▶ mutual exclusion protocols, barriers, reductions, ...
- ▶ **state**: a flag has two values, 0 and 1
- ▶ **atomic operations**
 - ▶ **raise**
 - ▶ can only be invoked when value is 0, otherwise **error**
 - ▶ sets value to 1
 - ▶ **lower**
 - ▶ blocks until value is 1, then sets value to 0 in one atomic step
 - ▶ no other thread can perform any operation on flag between check that value is 1 and set to 0

Interface for flags: `flag.h`

```
typedef ... flag_t;

/* Initializes the flag with the given value.  Must be called before
   the first time the flag is used. */
void flag_init(flag_t * f, _Bool val);

/* Destroys the flag */
void flag_destroy(flag_t * f);

/* Increments f atomically, and returns the result.  Notifies threads
   waiting for a change on f. An assertion is violated if f is 1 when
   this function is called. */
void flag_raise(flag_t * f);

/* Waits for f to be 1, then sets it to 0, all atomically. */
void flag_lower(flag_t * f);
```

Implementation of flags: `flags.h` and `flags.c`

```
typedef struct flag {
    _Bool val;
    pthread_mutex_t mutex;
    pthread_cond_t condition_var;
} flag_t;

void flag_init(flag_t * f, _Bool val) {
    f->val = val;
    pthread_mutex_init(&f->mutex, NULL);
    pthread_cond_init(&f->condition_var, NULL);
}

void flag_destroy(flag_t * f) {
    pthread_mutex_destroy(&f->mutex);
    pthread_cond_destroy(&f->condition_var);
}
```

Implementation of flags: raise and lower

```
void flag_raise(flag_t * f) {
    pthread_mutex_lock(&f->mutex);
    assert(!f->val);
    f->val = 1;
    pthread_cond_broadcast(&f->condition_var);
    pthread_mutex_unlock(&f->mutex);
}

void flag_lower(flag_t * f) {
    pthread_mutex_lock(&f->mutex);
    while (f->val == 0)
        pthread_cond_wait(&f->condition_var, &f->mutex);
    f->val = 0;
    pthread_mutex_unlock(&f->mutex);
}
```

Application of flags: barrier implementations

A very common pattern in multi-threaded programs:

```
while (true) {  
    compute something;  
    barrier();  
}
```

- ▶ `barrier()`: no thread can leave until every thread has entered
- ▶ thread 1 needs to read something produced by thread 2 in previous iteration
- ▶ how to construct a “barrier” for threads?
- ▶ many ways, using synchronization primitives we have already learned
- ▶ solutions differ in their performance characteristics
- ▶ desired characteristics of barriers:
 1. no one leaves until everyone enters
 2. no unnecessary delay: after last thread enters, everyone can leave without further delay
 3. **reuseable** : need to use the same barrier object over and over

A 2-thread barrier using flags

- ▶ two flags are used `f1` and `f2`
 - ▶ `f1` is used by Thread 1 to send a signal to Thread 2 saying “I have arrived at barrier”
 - ▶ `f2` is used by Thread 2 to send a signal to Thread 1 saying “I have arrived at barrier”
- ▶ Thread 1
 1. raises `f1`
 2. lowers `f2`
- ▶ Thread 2
 1. lowers `f1`
 2. raises `f2`

Is it a correct, re-useable barrier with no unnecessary delay?

See [2barrier.c](#).