

CISC 372: Parallel Computing

Threads, part 4: barrier and reduction implementations

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

Barrier implementations

Last lecture, we saw a simple two-thread barrier.

- ▶ two flags are used: `f1` and `f2`
 - ▶ `f1` is used by Thread 1 to send a signal to Thread 2 saying “I have arrived at barrier”
 - ▶ `f2` is used by Thread 2 to send a signal to Thread 1 saying “I have arrived at barrier”
- ▶ Thread 1
 1. raises `f1`
 2. lowers `f2`
- ▶ Thread 2
 1. lowers `f1`
 2. raises `f2`

Now let's generalize...

A counter barrier

- ▶ a shared counter keeps track of number of threads in barrier
- ▶ last thread to enter resets counter and signals all other threads to depart
- ▶ one flag for each thread is used to transmit the departure signal
- ▶ see `flag_barrier.c`

Analysis

- ▶ (+) only 1 flag per thread
- ▶ (-) **contention** on shared variable counter (and its lock)
- ▶ (-) $O(n)$ due to last thread's protocol

A coordinator barrier

- ▶ use an additional “**coordinator**” thread
 - ▶ a special thread that is not part of the “team”
 - ▶ it keeps track of who is in the barrier
 - ▶ the n regular threads are “**workers**”
- ▶ use two flags for each worker
 1. for worker to signal coordinator it has arrived
 2. for coordinator to signal worker it may leave
- ▶ worker barrier protocol
 1. worker signals coordinator “I have arrived”
 2. worker waits for departure signal from coordinator and lowers flag
- ▶ coordinator protocol
 1. loop over workers:
 - ▶ wait for and lower arrival flag for each
 2. loop over workers:
 - ▶ send departure signal to each



▶ see [coordinator_barrier.c](#)

Analysis of coordinator barrier

- ▶ (+) avoids memory contention
- ▶ (-) requires an extra thread
- ▶ (-) $O(n)$: execution time is proportional to n
 - ▶ what if $n = 10^6$?



Combining binary tree barrier

- ▶ combines actions of worker and coordinator
 - ▶ so each worker also coordinates
- ▶ organizes workers in tree
 - ▶ execution time proportional to $\log(n)$
- ▶ flow of signals
 - ▶ send arrive signals up the tree
 - ▶ send depart signals down the tree
- ▶ sequence of events
 - ▶ worker waits for all children to arrive
 - ▶ then tells parent it has arrived
 - ▶ when root learns that its children have arrived
 - ▶ it knows all procs have arrived
 - ▶ then root tells its children to depart
 - ▶ when a worker is told to depart, it tells its children to depart, ...



Combining binary tree barrier: `tree_barrier.c`

- ▶ protocol for leaf node L
 1. raise `arrive`[L]
 2. lower `depart`[L]
- ▶ protocol for interior node I
 1. lower `arrive`[*left*]
 2. lower `arrive`[*right*]
 3. raise `arrive`[I]
 4. lower `depart`[I]
 5. raise `depart`[*left*]
 6. raise `depart`[*right*]
- ▶ protocol for root node R
 1. lower `arrive`[*left*]
 2. lower `arrive`[*right*]
 3. raise `depart`[*left*]
 4. raise `depart`[*right*]



Analysis of combining tree barrier


- ▶ time is $O(\log(n))$
 - ▶ no loops 1.. n
 - ▶ each row in the tree can execute in parallel
- ▶ different procs execute different code
- ▶ leaf and root execute fewer instructions
 - ▶ could lead to inefficiency
- ▶ increases complexity

Symmetric barriers

- ▶ in **symmetric barriers**
 - ▶ all procs execute same code
- ▶ common structure: solutions are constructed from pairs of 2-process barriers
 - ▶ Thread 1
 1. raises **f1**
 2. lowers **f2**
 - ▶ Thread 2
 1. lowers **f1**
 2. raises **f2**

Symmetric barriers


- ▶ need to choose **interconnection scheme**
 - ▶ sequence of 2-process barriers executed by each proc
- ▶ Example: **butterfly barrier**
 - ▶ assume $n = 2^s$ is a power of 2



- ▶ butterfly: $O(\log(n))$ time and symmetric
- ▶ requires n to be power of 2

Dissemination Barrier: `dissem_barrier.c`

- ▶ butterfly requires n to be power of 2 or have exceptional code which breaks symmetry
- ▶ dissemination barrier works for any n
- ▶ uses **cyclic** order of threads
- ▶ two flags (**a** and **b**) for each thread, in each stage
- ▶ in stage i each thread
 - ▶ synchs with thread 2^i to the right using **a** and **b** flags of that thread
 - ▶ synchs with thread 2^i to the left using its **a** and **b** flags
- ▶ stages: $0 \leq i < \lceil \log_2 n \rceil$



Dissemination barrier: code

```
for (int stage=0, i=1; stage<nstages; stage++, i*=2) {  
    flag_raise(&bs->a[stage][(tid+i)%nthreads]);  
    flag_lower(&bs->a[stage][tid]);  
    flag_raise(&bs->b[stage][tid]);  
    flag_lower(&bs->b[stage][(tid+i)%nthreads]);  
}
```

Reductions

- ▶ each barrier algorithm can be extended to a reduction algorithm
 - ▶ for example, to sum elements of an array efficiently
- ▶ example: tree barrier
 - ▶ each node has an associated value
 - ▶ each node waits for arrival of its left and right child
 - ▶ then sets its value to sum of the value of the children
 - ▶ then alerts its parent, . . .
 - ▶ root gets the global sum and can assign it to a global variable
 - ▶ allows reduction without all threads contending for a single mutex

pthread_barrier

- ▶ Pthreads provides a barrier
 - ▶ but this is an **optional feature**
 - ▶ not supported on all platforms (even those supporting Pthreads)
 - ▶ for portable code: know how to write your own barrier
- ▶ `pthread_barrier_t` : type of a barrier object
- ▶ `pthread_barrier_init(...)`
 - ▶ `pthread_barrier_t *`
 - ▶ pointer to barrier object to initialize
 - ▶ `pthread_barrierattr_t *`
 - ▶ `unsigned int count`
 - ▶ number of threads that will participate in this barrier
- ▶ `pthread_barrier_destroy(pthread_barrier_t *)`
- ▶ `pthread_barrier_wait(pthread_barrier_t *)`