

CISC 372: Parallel Computing OpenMP

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

OpenMP Overview

- ▶ an API for shared memory, multi-threaded programming
- ▶ works with C, C++, or Fortran
- ▶ emphasizes **incremental** parallelization
 - ▶ start with the sequential program
 - ▶ add a little bit of parallelism at a time
 - ▶ see how it works, change it, add some more, ...
 - ▶ contrast with MPI “all or nothing” approach
- ▶ programmer inserts *directives* and function calls into sequential program
 - ▶ in C, a directive is a **pragma**
 - ▶ stands for **pragmatic information**
 - ▶ a general way to pass additional information to the compiler in a form not supported by the C language
 - ▶ a compiler that does not recognize a kind of pragma can just ignore it
 - ▶ `#pragma omp ...`
- ▶ the sequential program remains embedded in the OpenMP version
 - ▶ just ignore the pragmas
 - ▶ replace function calls with trivial implementations

Sequential Dot Product (Chapman et al., *Using OpenMP*)

```
#include<stdio.h>
int main() {
    double sum, a[256], b[256];
    int status, i, n=256;
    for (i = 0; i < n; i++) {
        a[i] = i * 0.5;
        b[i] = i * 2.0;
    }
    sum = 0;
    for (i = 0; i < n; i++) {
        sum = sum + a[i]*b[i];
    }
    printf("sum = %f \n", sum);
}
```

Dot Product in OpenMP (Chapman et al., *Using OpenMP*)

```
#include<stdio.h>
int main() {
    double sum, a[256], b[256];
    int status, i, n=256;
    for (i = 0; i < n; i++) {
        a[i] = i * 0.5;
        b[i] = i * 2.0;
    }
    sum = 0;
    #pragma omp parallel for reduction(+:sum)
    for (i = 0; i < n; i++) {
        sum = sum + a[i]*b[i];
    }
    printf("sum = %f \n", sum);
}
```

Basic Syntactic Concepts

- ▶ most directives are applied to the following **structured block** S
- ▶ S may be almost any kind of statement
 - ▶ a compound statement `{...}` (this is most common)
 - ▶ S must have single point of entry and single point of exit
 - ▶ may be a loop
 - ▶ if enclosing in curly braces would be a structured block
 - ▶ may be an *if* statement
 - ▶ if enclosing in curly braces would be a structured block
- ▶ in C, the directives have the form `#pragma omp ...`
- ▶ you can put non-newline white space before or after the `#`
- ▶ the directive terminates with the end of the logical line
- ▶ long directives can be spread over multiple physical lines by ending each physical line but the last with `\`
 - ▶ in C, these physical lines are merged into one logical line at a very early stage of compilation (before preprocessing)

Spreading a logical line over multiple physical lines

```
#pragma omp this is my really big long \  
    pragma that keeps going and going and \  
    going on and on and on and on and on \  
    and on and on  
for (i=0; i<n; i++) {  
    ...  
}
```

Beware: You cannot have any white space after the `\`. It must be the last character on the physical line.

Compiling and running an OpenMP program

- ▶ use gcc or clang, add flag `-fopenmp`; everything else the same

```
gcc -fopenmp -o dot dot.c
./dot
```
- ▶ without the flag `-fopenmp`
 - ▶ header file `omp.h` will not necessarily be found
 - ▶ pragmas will just be ignored; program will be sequential
- ▶ Apple users
 - ▶ for reasons that escape me, Apple's version of clang does not have OpenMP support
 - ▶ advice: install clang yourself using MacPorts
 - ▶ `sudo port install clang-9.0`, or later
 - ▶ then use `clang-mp-9.0`
- ▶ general: all compilers
 - ▶ the preprocessor object macro `_OPENMP` is defined iff you are running the compiler with OpenMP support
 - ▶ value is `yyymm`, where `yyyy` is the year of the Standard supported, and `mm` is the month
 - ▶ permits things like `#ifdef _OPENMP ... #else ...`

The `parallel` directive

```
#pragma omp parallel [clauses]  
S
```

- ▶ a program begins execution with one thread
- ▶ executing a **parallel** directive creates a **parallel region**
- ▶ when control enters the region, a **team** of threads is created
 - ▶ the team includes the original thread, known as the **master thread**
- ▶ all of the threads in the team execute the statement *S* concurrently
- ▶ *S* is typically a big compound statement
- ▶ additional directives inside *S* control how threads in the team behave
- ▶ at the end of *S* there is an implicit barrier
 - ▶ all threads join up at this point
 - ▶ all threads other than the master essentially disappear
- ▶ the master continues execution

hello1.c: parallel directive example

```
#include <stdio.h>

int main () {
    printf("I am the master.\n"); // just the master
#pragma omp parallel
    {
        printf("Hello, world.\n"); // all threads
    } /* end of parallel region */
    printf("Goodbye, world.\n"); // just the master
}
```

```
omp$ cc -fopenmp hello1.c
omp$ ./a.out
I am the master.
Hello, world.
Hello, world.
Goodbye, world.
omp$
```

Basic OMP functions

- ▶ need to `#include <omp.h>`
- ▶ `int omp_get_num_threads()`
 - ▶ returns the number of threads in the team in the current region
- ▶ `int omp_get_thread_num()`
 - ▶ returns the ID of the calling thread
 - ▶ threads within a team are numbered 0, 1, ...
 - ▶ master thread is always thread 0
- ▶ `omp_get_wtime()`
 - ▶ returns the wall clock time (like `MPI_Wtime`)