

CISC 372: Parallel Computing

OpenMP, Part 2

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

Private vs. shared variables

- ▶ if a variable is declared within the parallel region...
 - ▶ all threads have their own copy of that variable
- ▶ if a variable is declared before the parallel region and is visible in the region...
 - ▶ you have a choice: variable can be
 - ▶ **private**: all threads get their own private copy of the variable (in addition to the original), or
 - ▶ **shared**: one shared variable
- ▶ specify what you want by clauses of the form
 - ▶ `shared(u1,u2,...)`
 - ▶ `private(v1,v2,...)`
- ▶ some (obvious) points
 - ▶ the `u1,u2,...` and `v1,v2,...` must all be visible at this point
 - ▶ a variable cannot be both `shared` and `private`

The default clause

- ▶ sets the default for private vs. shared in the parallel region
- ▶ `default(none)`
 - ▶ no default
 - ▶ every variable used in parallel region must be explicitly listed in `shared` or `private`
- ▶ `default(shared)`
 - ▶ if not listed, the variable is shared
- ▶ there are rules specifying what happens if you don't have a default clause
 - ▶ but ignore them for now
 - ▶ explicitly declare every variable used in the region as either `private` or `shared`

hello2.c

```
#include <omp.h>
#include <stdio.h>

int main (int argc, char *argv[]) {
    int nthreads, tid;

    #pragma omp parallel private(nthreads, tid)
    {
        tid = omp_get_thread_num();
        printf("Hello World from thread = %d\n", tid);
        if (tid == 0) { // only master
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }
    } // end of parallel region
}
```

num_threads(): requesting the number of threads

- ▶ you can request that the team have a specified number of threads
- ▶ clause: `num_threads(expr)`
 - ▶ where `expr` is an expression which evaluates to a positive integer
- ▶ the runtime system **may** give you the requested number of threads
 - ▶ or it may give you fewer
- ▶ if you really need to know how many there are, ask
 - ▶ `int omp_get_num_threads()`

Initializing private variables

- ▶ suppose `x` is an integer variable declared before entering a parallel region
- ▶ value of `x` is 5 upon reaching the parallel region
- ▶ `x` is declared private
- ▶ when control enters the parallel region, what is the initial value of `x`?
- ▶ answer: undefined, `even on master thread`
- ▶ try it! see `initial.c`
- ▶ if you want the private `x` to be initialized with the value the original `x` had:
 - ▶ use `firstprivate`
- ▶ clause: `firstprivate(v1,v2,...)`
 - ▶ declares `v1,v2,...` to be not only private, but to be initialized with global value
 - ▶ see `initial2.c`

Constructs, regions, and modification of original variable of `private`

OpenMP 4.0 Sec. 1.2.2:

construct. *An OpenMP executable directive ... and the associated statement, loop or structured block, if any, not including the code in any called routines. That is, in the lexical extent of an executable directive.*

region. *All code encountered during a specific instance of the execution of a given construct or of an OpenMP library routine. A region includes any code in called routines as well as any implicit code introduced by the OpenMP implementation.*

Sec. 2.14.3.3, `private` clause:

*The value ... of the **original** list item will change only*

- ▶ *if accessed and modified via pointer,*
 - ▶ *if possibly accessed in the region but outside of the construct, [or]*
 - ▶ *as a side effect of directives or clauses[.]*
- ▶ **beware!** when you access a “private” variable outside of the construct
- ▶ you may be accessing the original copy; see `semiprivate.c`

Work-sharing

- ▶ you usually don't want all threads in the team to do the same thing
- ▶ you can code in branches on thread ID manually, but this is very tedious
- ▶ OpenMP provides more convenient, higher-level constructs
 - ▶ these are specified using directives within a parallel region
- ▶ one class of such constructs are the **work-sharing** constructs
 - ▶ these specify how work is to be divided up among members of the team
 - ▶ kinds of work-sharing constructs
 - ▶ **for** loops: distribute iterations to team members
 - ▶ **sections**: distribute independent code blocks (work units)
 - ▶ **single**: let only one thread execute a block

Worksharing constructs: `for` loops

► syntax

```
#pragma omp for [clauses]  
for (init-expr; var relop b; incr-expr)  
    body
```

► semantics

- each iteration is executed by exactly one thread in the team
- barrier at end of loop
- in general, everything else is unspecified
 - how the iterations are distributed among the team members
 - the order in which the iterations are executed
 - what happens concurrently

► syntactic restrictions on the `for` statement:

- `init-expr`: `var = expr`, integer type
- `relop` is one of: `<`, `<=`, `>`, `>=`
- `b` is a **loop-invariant** integer expression
- `incr-expr` has one of a few forms; see OpenMP 4.0 Standard, Section 2.6

Allowed forms for increment expression in `for` loop

- ▶ `++var`
- ▶ `var++`
- ▶ `--var`
- ▶ `var--`
- ▶ `var += incr`
- ▶ `var -= incr`
- ▶ `var = var + incr`
- ▶ `var = incr + var`
- ▶ `var = var - incr`

where `incr` is a **loop invariant integer expression**

- ▶ i.e., throughout one execution of the loop
 - ▶ `incr` will have the same value each time control reaches the top of the loop
- ▶ however `incr` could have different values in different loop executions

Loop invariant expressions

```
for (i=0; i<n; i++) {  
    /* no writes to n */  
}
```

Is **n** loop invariant? **Yes**

```
for (i=0; i<n; i++) {  
    for (j=0; j<i; j++) {  
        /* no writes to i,j,n */  
    }  
}
```

Is **i** invariant of inner loop? **Yes**

```
for (i=0; i<n; i++) {  
    int max = i;  
    for (j=0; j<max; j++) {  
        ...  
        if (a[j]>max) max = a[j];  
        ...  
    }  
}
```

Is **max** invariant of inner loop?
Probably not

Combining constructs

```
#pragma omp parallel  
#pragma omp for  
S
```

can be abbreviated

```
#pragma omp parallel for  
S
```

- ▶ other constructs can be abbreviated similarly
- ▶ this is useful when you have just one construct inside a parallel region
- ▶ clauses must be unambiguous
 - ▶ if a clause is applicable only to **parallel**, fine
 - ▶ if a clause is applicable only to **for**, fine
 - ▶ if a clause is applicable to **parallel** and **for**
 - ▶ if it has the same meaning for each (e.g., **shared**), no problem
 - ▶ otherwise, **undefined behavior**

Clauses for the for loop directive

- ▶ `private(v1,v2,...)`
 - ▶ make a shared variable private for the duration of the loop
- ▶ `firstprivate(v1,v2,...)`
 - ▶ make a variable private and initialize it in every thread
- ▶ `lastprivate(v1,v2,...)`
 - ▶ make a variable private and copy the final value of variable in the **last iteration** back to the shared variable at end
- ▶ `reduction(...)`
 - ▶ apply some associative and commutative operation (like **+**) across all iterations for some variable
- ▶ `ordered`: declares that an `ordered` construct may occur in loop body
- ▶ `schedule`: options to control how iterations are distributed to threads
- ▶ `nowait`: remove the barrier at the end of the loop

Clauses for the for loop directive, cont.

- ▶ `collapse(n)`: apply directive to next *n* loops in a **loop nest**
 - ▶ *n* is an expression that evaluates to a positive integer
 - ▶ iteration space of the *n* loops is collapsed into a single space
 - ▶ the iterations in the resulting space are distributed to threads
 - ▶ all initializers, incremeters, and conditions must be invariant under all loops
 - ▶ i.e., they must remain constant throughout the entire loop nest

Correct:

```
#pragma omp for collapse(2)
for (i=0; i<n; i++)
    for (j=0; j<m; j++)
        a[i][j] = 2*b[i][j];
```

Incorrect:

```
#pragma omp for collapse(2)
for (i=0; i<n; i++)
    for (j=i; j<m; j++)
        a[i][j] = 2*b[i][j];
```

Question 1

Assume **a** and **b** are disjoint arrays.

Can this loop be parallelized with an OpenMP **for** construct?

```
for (i=0; i<n && a[i]>0; i++)  
    b[i] = b[i] - a[i];
```

No (non-standard condition)

Question 2

Assume **a** and **b** are disjoint arrays.

Can this loop be parallelized with an OpenMP **for** construct?

```
for (i=1; i<n; i++)  
    b[i] = b[i] - a[i] + b[i-1] - a[i-1]
```

No (data race)

Question 3

Assume **a**, **b**, and **c** are disjoint arrays.

Can this loop be parallelized with an OpenMP **for** construct?

```
for (i=1; i<n; i++)  
    c[i] = b[i] - a[i] + b[i-1] - a[i-1]
```

Yes

Question 4

Can this loop be parallelized with an OpenMP `for` construct?

```
for (i=1; i<n; i+=k)
    c[i] = b[i] - a[i] + b[i-1] - a[i-1]
```

Yes