

CISC 372: Parallel Computing

OpenMP, Part 3

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

OpenMP worksharing directives

Recall:

- ▶ used to divide up work among threads
- ▶ kinds of work-sharing constructs
 - ▶ **for** loops: distribute iterations to team members
 - ▶ **sections**: distribute independent code blocks (work units)
 - ▶ **single**: let only one thread execute a block

We left off looking at different clauses that can be used with the **omp for** directive.

Reductions: `reduction(reduction-identifier : list)`

- ▶ this is another clause that can be added to an `omp for` directive
- ▶ performs an (approximately) associative and commutative operation across all threads
- ▶ each variable v in the list should be a shared variable
- ▶ v should be initialized before entering the loop
- ▶ effectively, a `private` copy of v is created
- ▶ each private v is initialized to the default initial value corresponding to the operation
 - ▶ 0 for $+$, 1 for $*$, etc.
- ▶ all operations in loop body take place on the private copies
- ▶ when a thread finishes its iterations:
 - ▶ it adds (or whatever the operation is) its private value back to the shared v
 - ▶ this happens `atomically` to prevent races

Reduction example: `reduce.c`

```
#include <stdio.h>
#include <omp.h>
#define n 10
int a[n], s=1000000;
int main() {
    printf("Start s = %d\n", s);
    #pragma omp parallel default(none) shared(a,s)
    {
        int tid = omp_get_thread_num();
        #pragma omp for
        for (int i=0; i<n; i++) a[i] = i;
        #pragma omp for reduction(+:s) schedule(static,1)
        for (int i=0; i<n; i++) {
            s+=a[i];
            printf("Local s on thread %d = %d\n", tid, s);
        }
    }
    printf("Final s = %d\n", s);
}
```

Reduction example: output

```
omp$ make reduce
cc -fopenmp -o reduce.exec reduce.c
./reduce.exec
Start s = 1000000
Local s on thread 0 = 0
Local s on thread 0 = 2
Local s on thread 0 = 6
Local s on thread 0 = 12
Local s on thread 0 = 20
Local s on thread 1 = 1
Local s on thread 1 = 4
Local s on thread 1 = 9
Local s on thread 1 = 16
Local s on thread 1 = 25
Final s = 1000045
omp$
```

Reduction operations

operation	operator	initial value
addition	+	0
multiplication	*	1
subtraction (?)	-	0
bitwise and	&	~0
bitwise or		0
bitwise exclusive or	^	0
logical and	&&	1
logical or		0

Controlling loop schedules: `schedule(static, chunk_size)`

- ▶ iterations are partitioned into chunks of size *chunk_size*
- ▶ chunks are distributed in round-robin order to threads
- ▶ last chunk may be smaller
- ▶ distribution is “**static**”: determined upon reaching the loop
- ▶ you can omit *chunk_size*
 - ▶ iteration space divided into chunks of approximately equal size
 - ▶ at most one chunk given to each thread

Controlling loop schedules: `schedule(dynamic, chunk_size)`

- ▶ iterations are partitioned into chunks of size *chunk_size*
- ▶ chunks are distributed to threads **as they request them**
 - ▶ similar to the “manager-worker” pattern
 - ▶ as soon as a thread completes its chunk, it asks for a new one
- ▶ last chunk may be smaller
- ▶ advantageous when time to execute an iteration varies in an unpredictable way
- ▶ distribution is “**dynamic**”: determined as loop executes

Controlling loop schedules: `schedule(guided, chunk_size)`

- ▶ this is a variation on *dynamic* in which the chunk size **decreases** as execution proceeds
- ▶ size of chunk proportional to number of unassigned iterations divided by number of threads
 - ▶ *chunk_size* is a **lower bound** on the size of a chunk
 - ▶ for *chunk_size* = 1, size of a chunk decreases to 1
 - ▶ for *chunk_size* = $k > 1$, all chunks other than last must contain at least k iterations
- ▶ motivation
 - ▶ there is overhead to the manager-worker protocol
 - ▶ bigger chunks \rightarrow less overhead, but greater probability of leaving a thread idle
 - ▶ compromise: increase granularity as iteration space gets smaller, when the chance of leaving a thread idle is greater

To wait or not to wait?

- ▶ use of `nowait` clause in `for` directive removes the implicit barrier at end of loop
- ▶ this can increase concurrency, and performance
- ▶ but **can also introduce bugs**
 - ▶ use with extreme caution
 - ▶ make sure it does not introduce **data races**

To wait or not to wait: wait1.c

```
int main () {
    double a[n], b[n];
    #pragma omp parallel default(none) shared(a,b)
    {
        #pragma omp for nowait
        for (int i=0; i<n; i++)
            a[i] = 2.0*i;
        #pragma omp for
        for (int i=0; i<n; i++)
            b[i] = 3.0*i;
    } /* end of parallel region */
    for (int i=0; i<n; i++) {
        if (a[i]!=2.0*i) { printf("Error at a[%d]: %f\n", i, a[i]); fflush(stdout); exit(1); }
        if (b[i]!=3.0*i) { printf("Error at b[%d]: %f\n", i, b[i]); fflush(stdout); exit(1); }
    }
    printf("Success\n");
}
```

- OK: the two loops can execute concurrently since they update distinct variables

To wait or not to wait: wait2.c

```
int main() {
    double a[n], b[n];
    #pragma omp parallel default(none) shared(a,b)
    {
        #pragma omp for nowait
        for (int i=0; i<n; i++) a[i] = 2.0*i;
        #pragma omp for
        for (int i=0; i<n; i++) b[i] = 2.0*a[n-i-1];
    } /* end of parallel region */
    for (int i=0; i<n; i++) {
        if (a[i] != 2.0*i) {
            printf("Error at a[%d]: %f\n", i, a[i]); fflush(stdout); exit(1);
        }
        if (b[i] != 2.0*(2.0*(n-i-1))) {
            printf("Error at b[%d]: %f\n", i, b[i]); fflush(stdout); exit(1);
        }
    }
    printf("Success 2\n");
}
```

► **NOT OK**: second loop reads variables assigned in the first loop. Run it, and then run [wait2_fix.c](#).

Worksharing constructs: sections

```
#pragma omp sections
{
    #pragma omp section
    ...
    #pragma omp section
    ...
    :
}
```

- ▶ specifies explicit code blocks which can execute in parallel
- ▶ each block (or **section**) is executed once, by exactly one thread
- ▶ a thread may execute several sections, or no sections
- ▶ in general: you cannot assume anything about how sections are distributed to threads
- ▶ barrier at end (unless overridden with **nowait**)

sections example: sections.c, part 1

```
#include <stdio.h>
#include <omp.h>
#include <limits.h>
#define N 20
typedef unsigned long ulong;

ulong sumUpTo(int n) {
    ulong s=0;
    for (int i=1; i<=n; i++) s+=i;
    return s;
}

ulong productUpTo(int n) {
    ulong p=1;
    for (int i=1; i<=n; i++) p*=i;
    return p;
}
```

sections example: sections.c, part 2

```
int main() {
#pragma omp parallel
{ /* begin parallel region */
    int tid = omp_get_thread_num();
    if (tid == 0) printf("Number of threads: %d\n", omp_get_num_threads());
#pragma omp sections
    { /* begin sections */
#pragma omp section
    {
        printf("Thread %d: sum to %d ..... %lu\n", tid, N, sumUpTo(N));
    }
#pragma omp section
    {
        printf("Thread %d: product to %d ..... %lu\n", tid, N, productUpTo(N));
    }
    } /* end of sections */
} /* end of parallel region */
}
```

Clauses allowed with `sections`

- ▶ `private(list)`
 - ▶ each section has its own private copy of variable
- ▶ `firstprivate(list)`
 - ▶ make private and initialize with shared variable value
- ▶ `lastprivate(list)`
 - ▶ value of private copy of variable in last section is copied to shared variable at end
- ▶ `reduction(reduction-identifier:list)`
 - ▶ reduction applied across all sections
- ▶ `nowait`
 - ▶ removes barrier at end

Worksharing constructs: `single`

```
#pragma omp single  
S
```

- ▶ indicates that you want only one thread in the team to execute S
 - ▶ you don't care which thread
- ▶ barrier at end (unless overridden with `nowait`)
- ▶ typical use: initialization of shared variable

Clauses:

- ▶ `private(list)`, `firstprivate(list)`, `nowait`: usual semantics
- ▶ `copyprivate(list)`
 - ▶ applies to private variables
 - ▶ copies final value of variable in the single thread to corresponding variables in all other threads
 - ▶ copy occurs at end, before threads leave the barrier