

CISC 372: Parallel Computing

CUDA, part 1

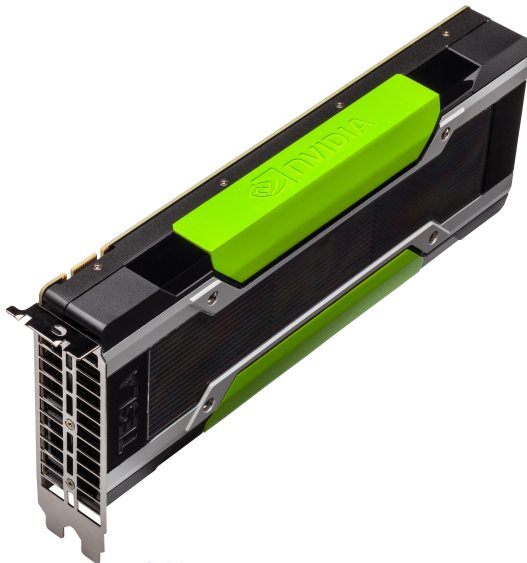
Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

CUDA Overview

- ▶ graphical processors (GPUs) are massively parallel machines
 - ▶ very different architecture from CPUs
 - ▶ multiple “streaming multiprocessors”
 - ▶ designed for doing large vector computations in parallel
 - ▶ not so good for algorithms with complex logic and branching
- ▶ ~ 2007: researchers realized that GPUs could be used for purposes other than graphics
 - ▶ birth of the **GPGPU**: general purpose GPU
- ▶ what was missing was a nice general purpose programming language targeting GPGPUs
- ▶ NVIDIA developed **CUDA-C** to fill this niche
 - ▶ based on C/C++
 - ▶ a few new language primitives; runtime library
 - ▶ for writing **heterogeneous** programs: mix of GPU and CPU code
 - ▶ **kernel**: a compute-intensive function to be run on the GPU
 - ▶ primitives for launching kernels, copying data between CPU and GPU
 - ▶ many algorithms can see orders of magnitude performance improvements over CPU

NVIDIA Tesla K80



K80 properties

```
Device name: Tesla K80
Compute capability: 3.7
Number of SMPs: 13
Max threads per block: 1024
Registers per block: 65536
Warp size: 32
Total global memory: 11996954624
Total constant memory: 65536
Shared memory per block: 49152
Memory Clock Rate (KHz): 2505000
Memory Bus Width (bits): 384
Peak Memory Bandwidth (GB/s): 240.480000
```

Price: \sim \$500

CUDA Background

- ▶ CUDA C is an extension of C for writing programs targeting NVIDIA's GPUs
- ▶ goal is to use GPUs for general purpose computing
- ▶ introduced in 2007, updated regularly
- ▶ some scientific problems can see enormous performance gains
- ▶ see <https://developer.nvidia.com/about-cuda>

References:

- ▶ the CUDA Programming Guide
 - ▶ in our repo under [docs/](#)
 - ▶ <https://docs.nvidia.com/cuda/cuda-c-programming-guide/>
- ▶ CUDA by Example
 - ▶ <https://developer.nvidia.com/cuda-example>
 - ▶ pay for the book, examples are free

CUDA Programming Model

- ▶ program execution starts as a single thread as usual
- ▶ a *kernel* function is a function declared with specifier `__global__`
- ▶ invoking a kernel function instantiates a **grid** executing on a **device**
- ▶ the grid consists of a collection of **blocks**
 - ▶ organized in a 1d, 2d, or 3d Cartesian geometry
- ▶ each block consists of a collection of **threads**
 - ▶ organized in a 1d, 2d, or 3d Cartesian geometry
- ▶ the blocks must be completely independent
 - ▶ they can execute concurrently or sequentially or anything in between
 - ▶ in any order
- ▶ the threads within a block execute concurrently and may coordinate
 - ▶ barriers
 - ▶ shared memory (shared by all threads in the block)

hello1.cu: Hello, world

```
#include <stdio.h>

__global__ void kernel(void) {
    printf("Hello from the GPU!\n");
}

int main (void) {
    kernel<<<1,1>>>(); // launch kernel with 1 block, 1 thread per block
    printf("Hello from the CPU!\n");
    cudaDeviceSynchronize(); // wait for kernel to return
}
```

- ▶ `__global__` indicates a function is a **kernel**: to be run on GPU
- ▶ `f<<<blocks,threadsPerBlock>>>(...)`
 - ▶ launch the kernel with `blocks` blocks and `threadsPerBlock` threads per block
 - ▶ returns immediately as kernel runs concurrently on GPU

Compiling and running a CUDA program on Beowulf

```
siegel@grendel:~/372/code/src/cuda/hello$ nvcc -o hello1.exec hello1.cu
siegel@grendel:~/372/code/src/cuda/hello$ srun -n 1 --gres=gpu:1 ./hello1.exec
srun: job 172804 queued and waiting for resources
srun: job 172804 has been allocated resources
Hello from the CPU!
Hello from the GPU!
siegel@grendel:~/372/code/src/cuda/hello$
```

- ▶ `nvcc`: similar to `cc`, different options
- ▶ `--gres=gpu:1` requests one GPU (in addition to the one CPU core)

hello2.cu: multiple blocks, threads per block

```
#include <stdio.h>

__global__ void kernel(void) {
    int bid = blockIdx.x; // block ID number
    int tid = threadIdx.x; // thread ID number (within its block)
    printf("Hello from block %d, thread %d of the GPU\n", bid, tid);
}

int main (void) {
    kernel<<<3,4>>>(); // 3 blocks, 4 threads per block
    printf("Hello, World\n");
    cudaDeviceSynchronize();
}
```

Output of hello2.cu

```
$ nvcc -o hello2.exec hello2.cu
$ srun --unbuffered -n 1 --gres=gpu:1 ./hello2.exec
srun: job 172815 queued and waiting for resources
srun: job 172815 has been allocated resources
Hello, World
Hello from block 0, thread 0 of the GPU
Hello from block 0, thread 1 of the GPU
Hello from block 0, thread 2 of the GPU
Hello from block 0, thread 3 of the GPU
Hello from block 2, thread 0 of the GPU
Hello from block 2, thread 1 of the GPU
Hello from block 2, thread 2 of the GPU
Hello from block 2, thread 3 of the GPU
Hello from block 1, thread 0 of the GPU
Hello from block 1, thread 1 of the GPU
Hello from block 1, thread 2 of the GPU
Hello from block 1, thread 3 of the GPU
```

Compiling and running CUDA programs on Bridges-2

Create a batch script like this:

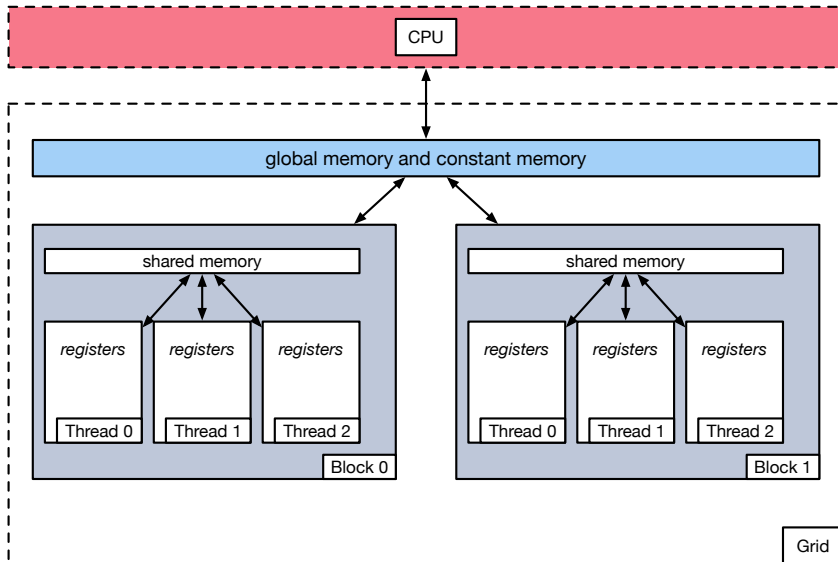
```
#!/bin/bash
#SBATCH -p GPU-shared
#SBATCH -t 00:01:00
#SBATCH -N 1
#SBATCH --gpus=1
# echo commands to stdout
set -x
./hello1.exec
```

- ▶ GPU partitions: [GPU-shared](#), [GPU-small](#), [GPU](#)
 - ▶ you will get charged for a full node (all CPUs and GPUs) unless you use [GPU-shared](#)
- ▶ each GPU node has multiple CPUs and GPUs
 - ▶ 8 NVIDIA Tesla V100-32GB SXM2 GPUs
 - ▶ 2 Intel Xeon Gold 6248 “Cascade Lake” CPUs: 40 cores total; 2.50–3.90 GHz
 - ▶ <https://portal.xsede.org/psc-bridges-2>: *GPU Nodes, GPU & GPU-shared partitions*

CUDA C Language Elements: Function type qualifiers

- ▶ `__device__`: executed on device, callable from device only
- ▶ `__global__`: **kernel**, executed on device, callable from host
 - ▶ function must return `void`
 - ▶ callable from device for **compute quality** $> 3.x$
 - ▶ calls must specify execution configuration
 - ▶ asynchronous
- ▶ `__host__`: executed on host, callable on host only
 - ▶ default
 - ▶ may be used with `__device__`
 - ▶ function is duplicated
 - ▶ use preprocessor macro `__CUDA_ARCH__` in body to include code that may be just for device or just for host version of function; if this is defined you are in CUDA version, else host version

CUDA memory hierarchy



Variable type qualifiers

- ▶ `__device__`
 - ▶ may be used in conjunction with `__constant__` or `__shared__`
 - ▶ when alone, variable resides in global memory space on device
 - ▶ has lifetime of application
 - ▶ is accessible from all threads within the grid
 - ▶ is accessible from host through library functions
 - ▶ may be used with `__managed__` to be directly referenced from host code
- ▶ `__constant__`
 - ▶ variable resides on device in **constant memory**
 - ▶ has lifetime of application
 - ▶ accessible from all threads in grid
 - ▶ accessible from host through library functions
- ▶ `__shared__`
 - ▶ variable resides in shared memory space of one block
 - ▶ has lifetime of the block
 - ▶ only accessible by all threads in that block
- ▶ nothing: thread-local variable